

Contents

I	Week 3: Classification and Logistic Regression	1
0.1	Logistic regression: binary classification	2
0.1.1	Hypothesis representation	2
0.1.2	Decision boundary	3
0.1.3	Cost function for a classification problem	3
0.1.4	Gradient Descent	4
0.2	Advanced Optimization	5
0.3	Multiclass classification	6
0.3.1	One-vs-all (one-vs-rest) classification	6
0.4	Underfitting(bias), Overfitting(variance)	7
0.4.1	Regularization	7
0.5	Assignments	10
0.5.1	Visualizing the data	10
0.5.2	Create the sigmoid() function	11
0.5.3	cost function and gradient	11
0.5.4	Plot the decision boundary	12
0.5.5	Predict label	13
0.5.6	Learning parameters using <i>fminunc()</i>	13
	Appendices	16
.1	Logistic Regression - University of Washington Coursera	17

Machine Learning

June 2, 2016

Part I

Week 3: Classification and Logistic Regression




Figure 1: Sigmoid (Logistic) function: $g(z) = \frac{1}{1+e^{-z}}$

0.1 Logistic regression: binary classification

Logistic regression is a **classification algorithm** where the value y take only a small number of discrete values.

- **binary classification problem:** y can only take 2 values: $y \in \{0, 1\}$. ('0' is called the negative class and '1' the positive class). Given $x^{(i)}$, the corresponding $y^{(i)}$ is also called **the label** (for the training example).
- **multi-class classification problem:** $y \in \{0, 1, 2, 3\}$ or more.

0.1.1 Hypothesis representation

The classifier must output value between 0 and 1: $0 \leq h_{\theta}(x) \leq 1$, which can be achieved using *sigmoid* function:

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

We interpret $h_{\theta}(x)$ as the **estimate of the probability of $y = 1$ on input x** .

$$P(y = 1|x; \theta) = h_{\theta}(x)$$

$$P(y = 0|x; \theta) = 1 - h_{\theta}(x)$$

$$P(y = 1|x; \theta) + P(y = 0|x; \theta) = 1$$

$P(y = 1|x; \theta)$ reads: probability that $y = 1$ given x parameterized by θ .

Exple: if for $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumor size} \end{bmatrix}$, $h_{\theta}(x) = 0.7$, it implies that there is 70% chance that the tumor of the patient is malignant ($y = 1$).

0.1.2 Decision boundary


$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-(\theta^T x)}} = P(y = 1|x; \theta)$$

From the sigmoid curve, one can see that :

Predict $y = 1$ if $h_{\theta}(x) = g(\theta^T x) \geq 0.5 \Rightarrow \text{when } \theta^T x \geq 0$

Predict $y = 0$ if $h_{\theta}(x) = g(\theta^T x) < 0.5 \Rightarrow \theta^T x < 0$

Exple: Let's assume a hypothesis function: $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$, with $\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$



- Predict $y = 1$ if $\theta^T x = -3 + x_1 + x_2 \geq 0 \Rightarrow (x_1 + x_2) \geq 3$
- Predict $y = 0$ if $\theta^T x = -3 + x_1 + x_2 < 0 \Rightarrow (x_1 + x_2) < 3$

The decision boundary is defined by $h_{\theta}(x) = 0.5$ ($x_1 + x_2 = 3$)

0.1.3 Cost function for a classification problem

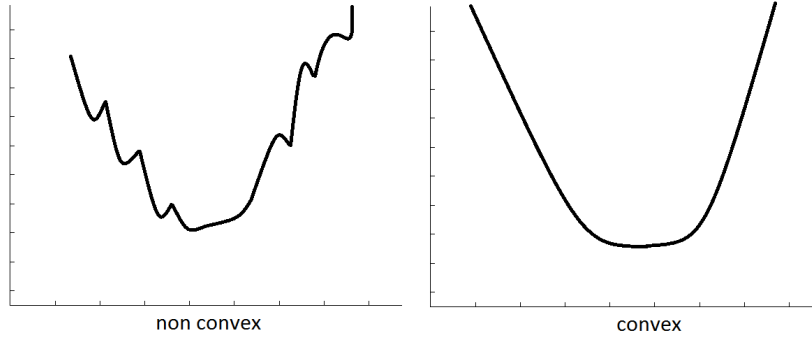
- Training set: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), (x^{(3)}, y^{(3)}), \dots, (x^{(m)}, y^{(m)})\}$ with m examples
- feature vector: $x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \in \mathbb{R}^{n+1}$ with $x_0 = 1$, and $y \in \{0, 1\}$
- The hypothesis: $h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$

For linear regression, we defined the cost function $J(\theta)$:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} (h_{\theta}(x^{(i)}) - y^{(i)})^2 = \frac{1}{m} \sum_{i=1}^m \text{Cost} [h_{\theta}(x^{(i)}), y^{(i)}]$$


where "Cost":

$$\text{Cost}(h_{\theta}(x), y) = \frac{1}{2} (h_{\theta}(x) - y)^2$$



For logistic regression, we use a different form of the function "Cost" and $J(\theta)$, so that $J(\theta)$ is convex, and gradient descent can converge:

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\text{Log}(h_{\theta}(x)) & \text{if } y = 1 \\ -\text{Log}(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$



The function 'Cost' can take a compact form:

$$\text{Cost}(h_{\theta}(x), y) = -y \text{Log}[h_{\theta}(x)] - (1 - y) \text{Log}[1 - h_{\theta}(x)]$$

If $y = 1$, $\text{Cost}(h_{\theta}(x), y) = -\text{Log}(h_{\theta}(x))$ like in granular form of the equation. Similarly for $y = 0$

Hence, $J(\theta)$ becomes:

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left(\sum_{i=1}^m y^{(i)} \text{Log}(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \text{Log}[1 - h_{\theta}(x^{(i)})] \right) \end{aligned}$$

0.1.4 Gradient Descent

1. Make prediction given $x \Rightarrow$ output: $h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}} = h_{\theta}(x) = P(y = 1|x; \theta)$
2. Calculate the cost function $J(\theta)$

3. Similarly to Batch Gradient Descent, we need $\min_{\theta} J(\theta)$.

Repeat until convergence {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

with

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(update all θ_j simultaneously: $j = 0$ to n)
}

Note that feature scaling can also be applied to logistic regression

0.2 Advanced Optimization

There are several optimization algorithms:

- Gradient Descent
 - conjugate gradient
 - BFGS
 - L – BFGS
- } No need to manually pick α . Faster than Gradient descent. But, they are also more complex.

(*advanced Numerical Computing *)

In MatLab, **fminunc()** (function minimization unconstrained) is a built-in advanced optimization function (`>> help fminunc`). In Python, the package **scipy.optimize** includes some optimization algorithms:

Here is the procedure for using **fminunc()**:

1. Step1: Generate a function **costFunction** that outputs 2 arguments

- *jVal*: the value of $J(\theta)$
- *gradient*: a vector with the value of the gradients

```
function [jVal, gradient] = costFunction(theta)
    jVal = [' code to compute J(theta) '];
    gradient = zeros(n+1, 1) #create a zero vector
    #and fill with gradient values
    gradient(1) = [' code to compute (d J(theta)/d theta_0) '];
    gradient(2) = [' code to compute (d J(theta)/d theta_1) '];
    .....
    .....
    gradient(n+1) = [' code to compute (d J(theta)/d theta_(n)) '];
```

2. Step2: Set the options

```
options = optimset('GradObj', 'on', 'MaxIter', '100');
```

- '100' is the maximum number of iterations
- 'GradObj', 'on' : tells **fminunc()** that our function returns both the cost and the gradient

Initial guess for theta

```
initialtheta = zeros(n+1,1)
```

3. Step3: run **fminunc()**

```
[optTheta, functionVal, exitFlag] = fminunc(@CostFunction, initialTheta, options)
```

'@CostFunction' is a pointer to the function **CostFunction()**.


4. Setp 4: **fminunc()** outputs 3 parameters values:

- **optTheta** = optimum value of θ as a vector
- **functionVal** (≈ 0) is the costFuntion value at optimum
- **exitFlag** (=1) shows convergence status.


0.3 Multiclass classification

Example of multiclass classification: email foldering/tagging

	work	friends	family	hobby
	$y = 1$	$y = 2$	$y = 3$	$y = 4$
or \rightarrow	$y = 0$	$y = 1$	$y = 2$	$y = 4$



Binary Classification



Multiclass classification

0.3.1 One-vs-all (one-vs-rest) classification



Transform a multiclass (k) problem into several (k) binary classes problems

Principle: Train a logistic regression classifier $h_{\theta}^{(i)}(x)$ for each class i to predict the probability that $y = i$.


On a new input x , to make a prediction, pick the class i that maximizes $h_{\theta}^{(i)}(x) \Rightarrow \max_i h_{\theta}^{(i)}(x)$.

Run all three classifier on input x and then choose classifier given the larger value ($\max_i h_{\theta}^{(i)}(x)$).


$$h_{\theta}^{(i)}(x) = P(y = i|x; \theta) \text{ with } i = 1, 2, 3$$



Classifier $h_{\theta}^{(1)}(x)$ - class 1 is positive class




Classifier $h_{\theta}^{(2)}(x)$ - class 2 is positive class




Classifier $h_{\theta}^{(3)}(x)$ - class 3 is positive class

0.4 Underfitting(bias), Overfitting(variance)




$$h_{\theta}(x) = \theta_0 + \theta_1 x$$

Underfitting/High Bias. There is a high bias of the model to fit the data to a linear despite the trend.



$$h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2$$



$$h_{\theta}(x) = \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

Overfitting/High Variance. There are too many features. The algorithm makes accurate predictions in the training example ($J(\theta) = 0$), but fails to generalize to new examples

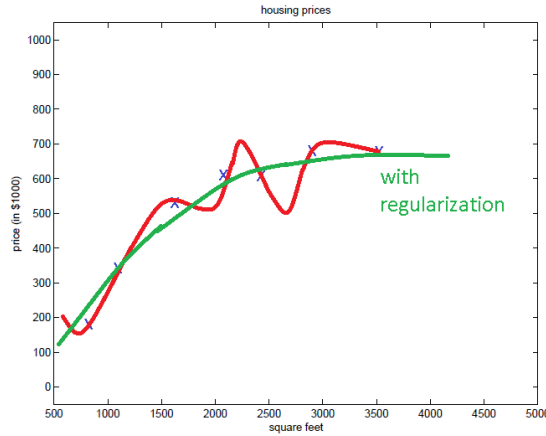
Bias/Variance also apply to logistic regression.

How to adress overfitting :

- Option1:
 - Reduce the number of features (select the more important features manually)
 - Model selection algorithm
- Option2: Regularization \Rightarrow keep all the features but reduce magnitude of θ_j

0.4.1 Regularization

The goal is to reduce the values/amplitudes of the parameters θ_j (with $j = 1, 2, \dots, n$), to mitigate overfitting (overconfidence). Often overfitting is associated with very large coefficients θ



- There are several options for the regularization term

– sum of squares (L_2 norm):

$$\| \theta \|_2 = \sum_{j=1}^n \theta_j^2 \quad (1)$$

– sum of absolute values (L_1 norm):

$$\| \theta \|_1 = \sum_{j=1}^n |\theta_j| \quad (2)$$

- Features: $\begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_n \end{bmatrix}$

- Parameters: $\begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix}$

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

λ is the regularization parameter. By convention, θ_0 is excluded from the regularization.

If λ is **too large** ($\lambda = 10^{10}$), algorithm would result in **underfitting** (fails to fit even the training set): $(\theta_1, \theta_2, \theta_3 \dots)$ parameters would be too much penalized $\approx 0 \Rightarrow \theta_1 x^{(1)} \approx \theta_2 x^{(2)} \dots \approx 0$ and hence $h_{\theta}(x) = \theta_0$.

Regularized linear regression

- Cost function

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

(θ_0 is excluded from the regularization term: $j = 1 \dots n$).

- Gradient descent: $\min_{\theta} J(\theta)$ with $h_{\theta}(x) = \theta^T x$

Repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

.....

.....

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j$$

for $j = 1, 2, 3 \dots n$

}

The equation of θ_j can be simplified:

$$\theta_j := \theta_j \left(1 - \alpha \frac{\lambda}{m} \right) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

Normal equation

With $X \in \mathbb{R}^{m \times (n+1)}$ and $y \in \mathbb{R}^m$

without regularization : $\theta = (X^T X)^{-1} X^T y$

$$\text{with regularization : } \theta = \left(X^T X + \lambda \begin{bmatrix} 0 & 0 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & 0 & \dots & 0 \\ 0 & 0 & 1 & 0 & \dots & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & \dots & 1 \end{bmatrix} \right)^{-1} X^T y$$

The matrix of (zeros and ones) is a $(n+1) \times (n+1)$.

Non-invertibility:

Suppose $m \leq n$ (# exles \leq #features), then $(X^T X)$ will be singular/non-invertible/degenerate (although MatLab can still provide a pseudo-inverse with '***pinv()***').

If $\lambda > 0$, then $(X^T X + \lambda M)$ -where M is the special matrix- is invertible: regularization makes the matrix invertible.

Regularized logistic Regression

- Cost function $J(\theta)$:

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \text{Log}(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \text{Log}(1 - h_{\theta}(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2$$

- Gradient descent: $h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$

Repeat until convergence {

$$\begin{aligned}\theta_0 &:= \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)} \\ \theta_j &:= \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \\ &\dots j = 1, 2, 3 \dots n\end{aligned}$$

}

With Advanced Optimization:

```
function [jVal, gradient] = costFunction(theta)
    jVal = [' code to compute J(theta) '];
    gradient = zeros(n+1, 1) #create a zero vector
    #and fill with gradient values
    gradient(1) = [' code to compute (d J(theta)/d theta_0) '];
    gradient(2) = [' code to compute (d J(theta)/d theta_1) '];
    .....
    .....
    gradient(n+1) = [' code to compute (d J(theta)/d theta_(n))      '];
```

where:

$$\begin{aligned}J(\theta) &= \left[-\frac{1}{m} \sum_{i=1}^m y^{(i)} \text{Log}(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \text{Log}(1 - h_{\theta}(x^{(i)})) \right] + \frac{\lambda}{2m} \sum_{j=1}^n \theta_j^2 \\ \frac{\partial}{\partial \theta_0} J(\theta) &= \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)} \\ \frac{\partial}{\partial \theta_1} J(\theta) &= \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)} \right] + \frac{\lambda}{m} \theta_1\end{aligned}$$

0.5 Assignments

0.5.1 Visualizing the data

```
function plotData(X, y)
    figure; # Create New Figure
    hold on;
    # X is the array [x0,x1,x2] of the m training examples
    # find indexes of training examples where y=1 (positive)
    pos = find(y==1);
    # find indexes of training examples where y=0 (negative)
```

```

neg = find(y==0);

m = length(y); #number of training examples
# k+ = black plus, ko=black circles
# MarkerSize= 7
# Linewidth=2
#MarkerFaceColor=yellow (filled circles)
plot(X(pos,1), X(pos,2), 'k+', 'LineWidth', 2, 'MarkerSize', 7);
plot(X(neg,1), X(neg,2), 'ko', 'MarkerFaceColor', 'y', 'MarkerSize', 7);
end

```

0.5.2 Create the sigmoid() function

```

function g = sigmoid(z)
    # return g the sigmoid of z
    # z can be a vector, a matrix or a scalar
    g = zeros(size(z));
    #Take the exponential of z and add 1
    # 1./(...) do an inverse element wise.
    g = 1./(1 + exp(-z))
end

```

0.5.3 cost function and gradient

The cost function in logistic regression is defined by:

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m [-y^{(i)} \text{Log}(h_{\theta}(x^{(i)})) - (1 - y^{(i)}) \text{Log}(1 - h_{\theta}(x^{(i)}))]$$

and the gradient (a vector of same length as θ where $j = 0, \dots, n$).

$$\frac{\partial J(\theta)}{\partial \theta} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

```

function [J, grad] = costFunction(theta, X, y)
    m = length(y); # number of training examples
    J = 0; #initial value of cost function
    grad = zeros(size(theta)); #initial matrix for gradient
    #co is the item inside the sum of J
    #h(x) = sigmoid(X*theta)
    co = (-y.*log(sigmoid(X*theta)))-(1-y).*log(1-sigmoid(X*theta));
    J = 1/m * sum(co) # J is a scalar
    #grad is a vector matrix with X'=transpose(X)
    grad = 1/m*X'*(sigmoid(X*theta) - y)
end

```

0.5.4 Plot the decision boundary

```
function plotDecisionBoundary(theta, X, y)
%PLOTDECISIONBOUNDARY Plots the data points X and y into a new figure with
%the decision boundary defined by theta
% PLOTDECISIONBOUNDARY(theta, X,y) plots the data points with + for the
% positive examples and o for the negative examples. X is assumed to be
% a either
% 1) Mx3 matrix, where the first column is an all-ones column for the
% intercept.
% 2) MxN, N>3 matrix, where the first column is all-ones

% Plot Data
plotData(X(:,2:3), y);
hold on

if size(X, 2) <= 3
    % Only need 2 points to define a line, so choose two endpoints
    plot_x = [min(X(:,2))-2, max(X(:,2))+2];

    % Calculate the decision boundary line
    plot_y = (-1./theta(3)).*(theta(2).*plot_x + theta(1));

    % Plot, and adjust axes for better viewing
    plot(plot_x, plot_y)

    % Legend, specific for the exercise
    legend('Admitted', 'Not admitted', 'Decision Boundary')
    axis([30, 100, 30, 100])
else
    % Here is the grid range
    u = linspace(-1, 1.5, 50);
    v = linspace(-1, 1.5, 50);

    z = zeros(length(u), length(v));
    % Evaluate z = theta*x over the grid
    for i = 1:length(u)
        for j = 1:length(v)
            z(i,j) = mapFeature(u(i), v(j))*theta;
        end
    end
    z = z'; % important to transpose z before calling contour

    % Plot z = 0
    % Notice you need to specify the range [0, 0]
    contour(u, v, z, [0, 0], 'LineWidth', 2)
end
hold off

end
```

0.5.5 Predict label

This function will predict the label (0 or 1) of based on the learned logistic regression parameters theta.

```
function p = predict(theta, X)
m = size(X, 1); # Number of training examples
p = zeros(m, 1);

# calculate sigmoid for all examples p is a vector matrix of size mx1
p = sigmoid(X*theta);
for student = 1:m
    if p(student) >= 0.5
        p(student)=1;
    else
        p(student)=0;
    end
end
end
```

0.5.6 Learning parameters using *fminunc()*

```
#initialization
#clear the variables, close all figs, clc=clear command window
clear ; close all; clc

# Load Data : col1=exam1, col2=exam2, col3=label(1 or 0)
data = load('ex2data1.txt');
X = data(:, [1, 2]); y = data(:, 3);

#===== Part 1: Plotting =====
fprintf(['Plotting data with + indicating (y = 1) examples ...\n' ...
        'and o indicating (y = 0) examples.\n']);
plotData(X, y);
#retains plots in the current axes so that
#new plots added to the axes do not delete existing plots
hold on;
# Labels and Legend
xlabel('Exam 1 score')
ylabel('Exam 2 score')
legend('Admitted', 'Not admitted')
hold off;

#===== Part 2: Compute Cost and Gradient =====
# Setup the data matrix appropriately, and add ones for the intercept term
[m, n] = size(X);
#Add intercept term to x and X_test
X = [ones(m, 1) X];

# Initialize fitting parameters
initial_theta = zeros(n + 1, 1);
```

```

# Compute and display initial cost and gradient
[cost, grad] = costFunction(initial_theta, X, y);

fprintf('Cost at initial theta (zeros): %f\n', cost);
fprintf('Gradient at initial theta (zeros): \n');
fprintf(' %f \n', grad);

#===== Part 3: Optimizing using fminunc =====
# Set options for fminunc
options = optimset('GradObj', 'on', 'MaxIter', 100);

# Run fminunc to obtain the optimal theta
#@t creates a function with argument t which calls the costFunction
[theta, cost, exit_flag] = fminunc(@(t)(costFunction(t, X, y)), ...
initial_theta, options);

% Print theta to screen
fprintf('Cost at theta found by fminunc: %f\n', cost);
fprintf('theta: \n');
fprintf(' %f \n', theta);

# Plot Boundary
plotDecisionBoundary(theta, X, y);
hold on;
xlabel('Exam 1 score')
ylabel('Exam 2 score')
legend('Admitted', 'Not admitted')

#===== Part 4: Predict and Accuracies: predict() =====
# After learning the parameters, you'll like to use it to predict the outcomes
# on unseen data. In this part, you will use the logistic regression model
# to predict the probability that a student with score 45 on exam 1 and
# score 85 on exam 2 will be admitted.

prob = sigmoid([1 45 85] * theta);
fprintf(['For a student with scores 45 and 85, we predict an admission ' ...
        'probability of %f\n\n'], prob);
# Compute accuracy on our training set
p = predict(theta, X);

fprintf('Train Accuracy: %f\n', mean(double(p == y)) * 100);

```




Figure 2: Training data with precision boundary

Appendices

.1 Logistic Regression - University of Washington Coursera

We define $l(\theta)$ as the likelihood of θ : it is a measure of the quality of the fit of the model to the training examples. Maximizing the likelihood over all possible θ will ensure to give the best classifier. For example, for 1 data, we pick a θ value that maximizes $p(y = +1|x, \theta) = \frac{1}{1+e^{-\theta^T x}}$.

- For m examples, the likelihood is the product of the probability (assuming that the examples are independent of each other):

$$l(\theta) = \prod_{i=1}^m p(y^{(i)}|x^{(i)}, \theta)$$

In order to simplify the mathematics, we use $\text{Log}l(\theta)$ (the Log does not change the maxima), hence:

$$\text{Log}(l(\theta)) = \sum_{i=1}^m \text{Log}(p(y^{(i)}|x^{(i)}, \theta))$$

We introduce the indicator function:

$$\mathbb{1}(y^{(1)} = +1) \begin{cases} +1 & \text{if } y^{(1)} \\ -1 & \text{if } y^{(1)} = -1 \end{cases} \quad (3)$$

For one example:

$$\begin{aligned} \text{Log}(l(\theta)) &= \mathbb{1}(y^{(1)} = +1)P(y^{(1)} = +1|x^{(i)}, \theta) + \mathbb{1}(y^{(1)} = -1)P(y^{(1)} = -1|x^{(i)}, \theta) \\ &= \mathbb{1}(y^{(1)} = +1)P(y^{(1)} = +1|x^{(i)}, \theta) + \mathbb{1}(y^{(1)} = -1)(1 - P(y^{(1)} = +1|x^{(i)}, \theta)) \\ &= \mathbb{1}(y^{(1)} = +1)P(y^{(1)} = +1|x^{(i)}, \theta) + \mathbb{1}(y^{(1)} = -1)\text{Log}\left(1 - \frac{1}{1 + e^{-\theta^T x}}\right) \\ &= \mathbb{1}(y^{(1)} = +1)P(y^{(1)} = +1|x^{(i)}, \theta) + \mathbb{1}(y^{(1)} = -1)\text{Log}\left(\frac{1 + e^{-\theta^T x} - 1}{1 + e^{-\theta^T x}}\right) \\ &= \mathbb{1}(y^{(1)} = +1)P(y^{(1)} = +1|x^{(i)}, \theta) + \mathbb{1}(y^{(1)} = -1)\text{Log}\left(\frac{e^{-\theta^T x}}{1 + e^{-\theta^T x}}\right) \end{aligned} \quad (4)$$

- Gradient:

$$\begin{aligned} \frac{\partial \text{Log}}{\partial \theta_j} &= -(1 - \mathbb{1}(y^{(1)} = +1))\frac{\partial \theta^T x}{\partial \theta_j} - \frac{\partial}{\partial \theta_j} \text{Log}\left(1 - e^{-\theta^T x}\right) \\ &= -(1 - \mathbb{1}(y^{(1)} = +1))x_j - \frac{-\theta_j e^{-\theta^T x}}{1 - e^{-\theta^T x}} \\ &= -(1 - \mathbb{1}(y^{(1)} = +1))x_j + x_j P(y^{(1)} = -1|x, \theta) \\ &= x_j (\mathbb{1}(y^{(1)} = +1) - P(y^{(1)} = +1|x, \theta)) \end{aligned} \quad (5)$$

For all examples m :

$$\frac{\partial \text{Log}}{\partial \theta_j} = \sum_{i=1}^N x_j (\mathbb{1}(y^{(i)} = +1) - P(y^{(i)} = +1|x, \theta)) \quad (6)$$